

Armeria

A Microservice Framework
Well-suited Everywhere

Trustin Lee, LINE
Oct 2019



@armeria_project



line/armeria

A microservice framework, *again*?



@armeria_project



online/armeria

Yeah, but for good reasons!

- Simple & User-friendly
- Asynchronous & Reactive
- 1st-class RPC support
 - ... with better-than-upstream experience
- Unopinionated integration & migration
- Less points of failure



@armeria_project



Qline/armeria

How simple is it, then?



@armeria_project



gline/armeria

Hello, world!

```
Server server = Server.builder()
    .http(8080)
    .https(8443)
    .tlsSelfSigned()
    .haproxy(8080)
    .service("/hello/:name",
        (ctx, req) -> HttpResponse.of("Hello, %s!",
            ctx.pathParam("name")))
    .build();
server.start();
```

Protocol auto-detection at 8080



@armeria_project



online/armeria

Hello, world – Annotated

```
Server server = Server.builder()
    .http(8080)
    .annotatedService(new Object() {
        @Get("/hello/:name")
        public String hello(@Param String name) {
            return String.format("Hello, %s!", name);
        }
    })
    .build();
server.start();
```

- Full example:

<https://github.com/line/armeria-examples/tree/master/annotated-http-service>



@armeria_project



Line/armeria



```
Server server = Server.builder()
    .http(8080)
    .service(GrpcService.builder()
        .addService(new GrpcHelloService())
        .build())
    .build();

class GrpcHelloService
    extends HelloServiceGrpc.HelloServiceImplBase {
    ...
}
```

- Full example:

<https://github.com/line/armeria-examples/tree/master/grpc-service>



@armeria_project



Line/armeria

Thrift

```
Server server = Server.builder()
    .http(8080)
    .service("/hello",
        THttpService.of(new ThriftHelloService()))
    .build();

class ThriftHelloService implements HelloService.AsyncIface {
    ...
}
```

Mix & Match!

```
Server server = Server.builder()
    .http(8080)
    .service("/hello/rest",
        (ctx, req) -> HttpResponse.of("Hello, world!"))
    .service("/hello/thrift",
        THttpService.of(new ThriftHelloService()))
    .service(GrpcService.builder()
        .addService(new GrpcHelloService())
        .build())
    .build();
```



@armeria_project



Qline/armeria

Why going asynchronous & reactive?

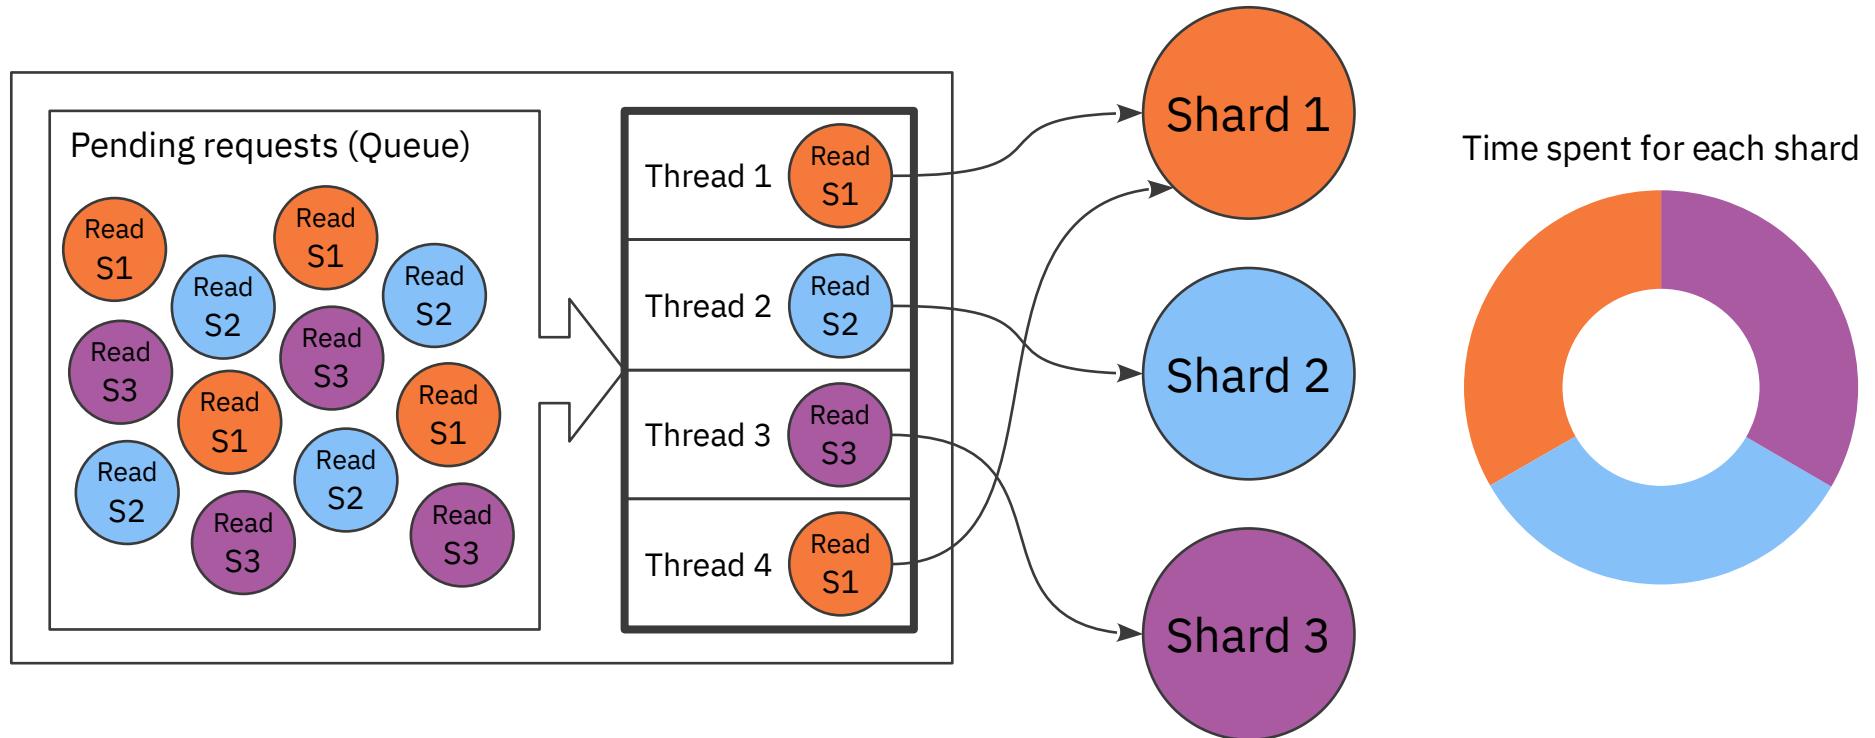


@armeria_project



online/armeria

One fine day of a synchronous microservice

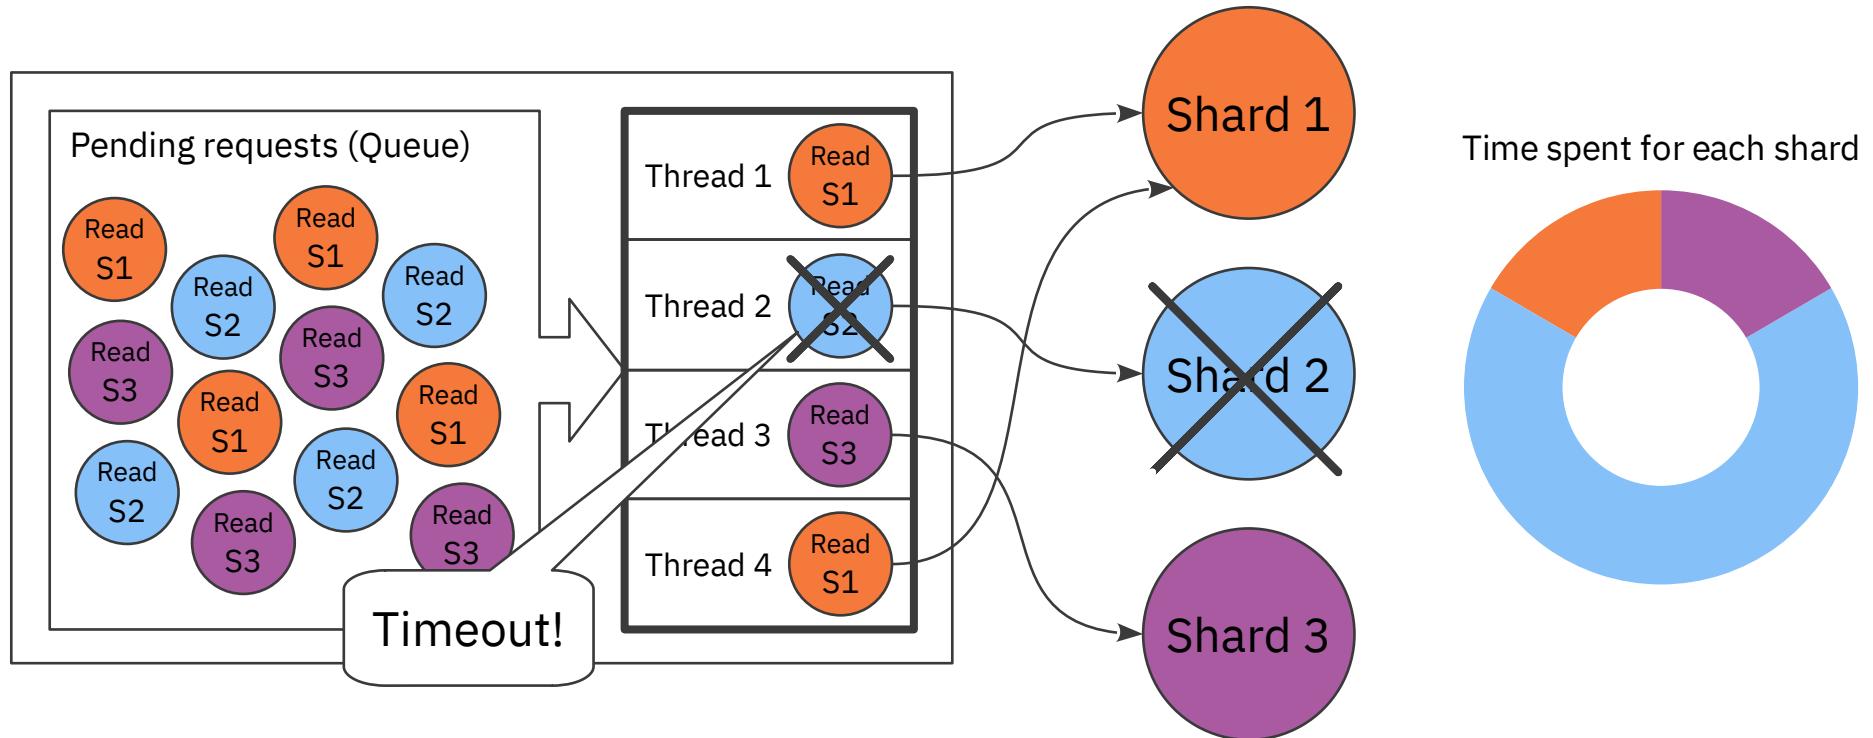


@armeria_project



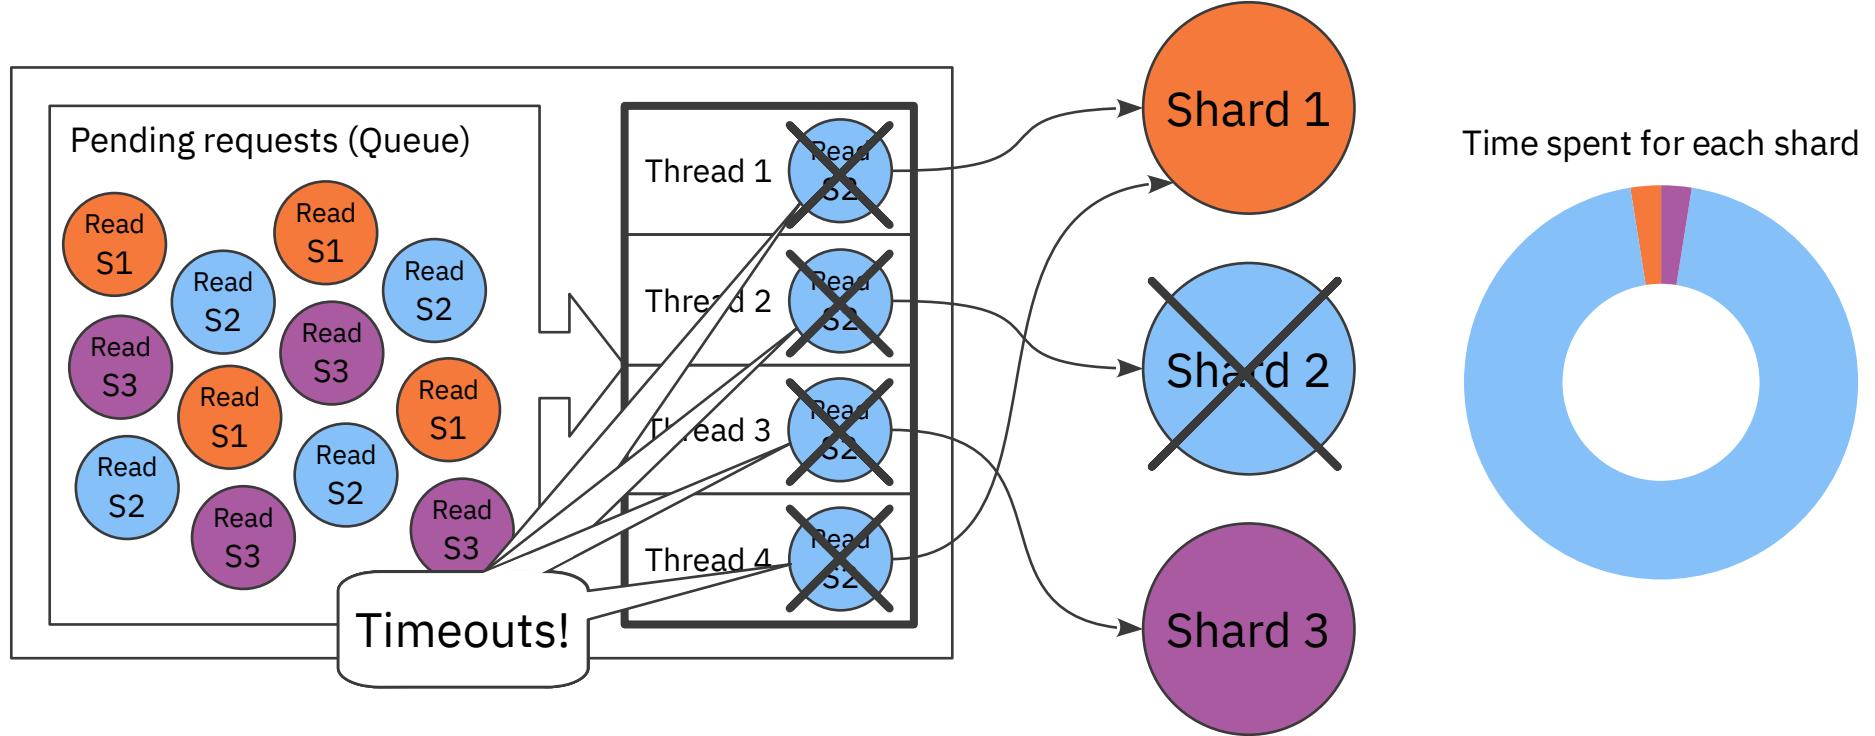
online/armeria

Shard 2 ruins the fine day...



Shard 1 & 3: Why are no requests coming? 🤔

Workers: We're busy waiting for Shard 2.

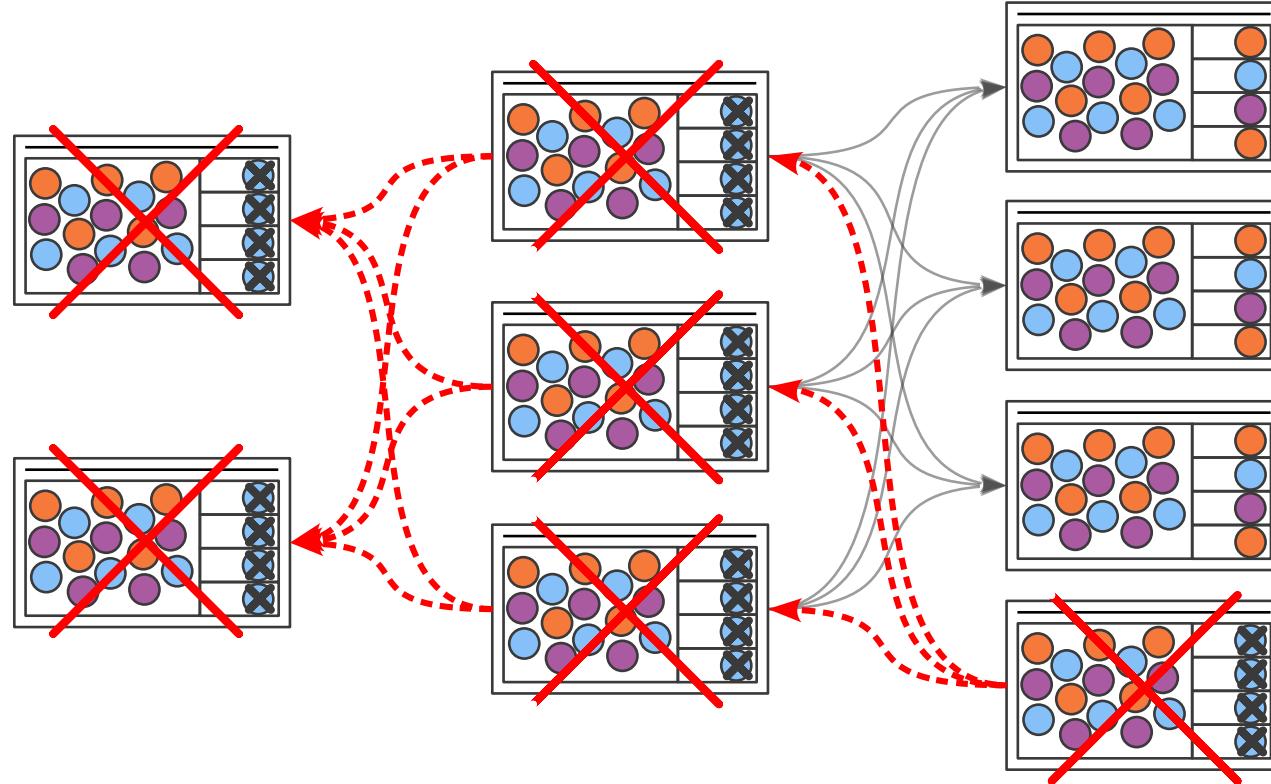


@armeria_project



online/armeria

... propagating everywhere!



@armeria_project



online/armeria

How can we solve this?

- Add more CPUs?
 - They are very idle.
- Add more threads?
 - They will all get stuck with Shard 2 in no time.
 - Waste of CPU cycles & memory – context switches & call stack
- Result:
 - Fragile system that falls apart even on a tiny backend failure
 - Inefficient system that takes more memory and CPU



@armeria_project



online/armeria

How can we solve this? (cont'd)

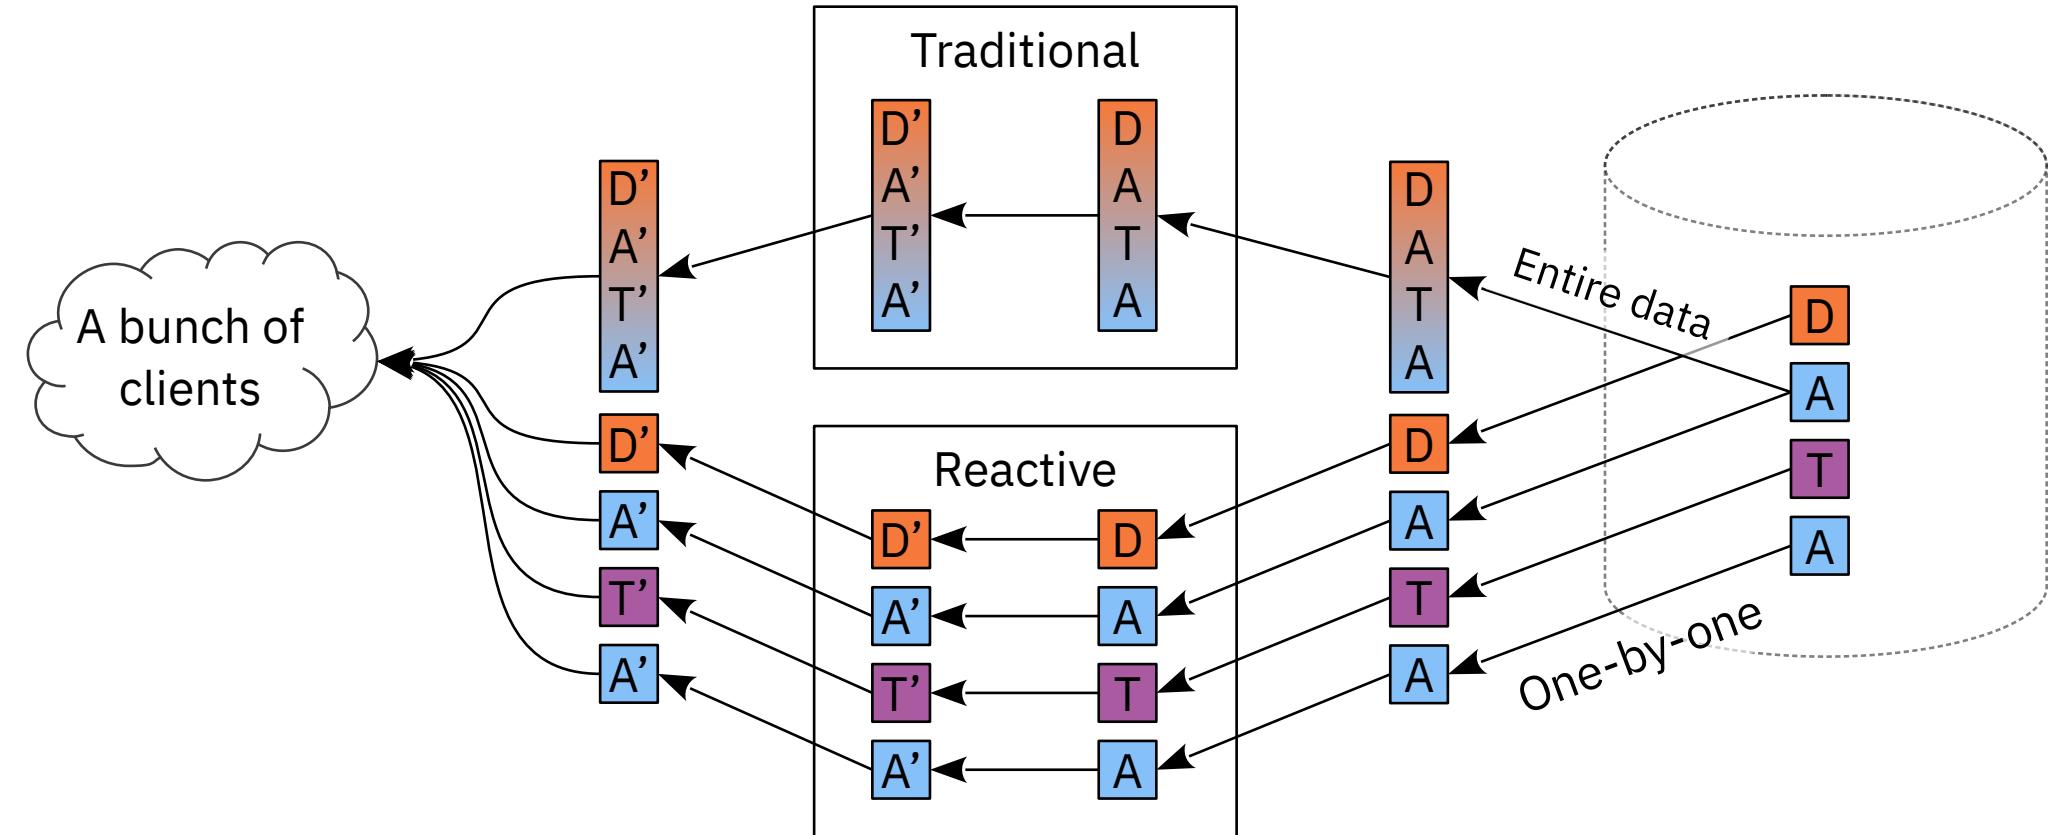
- Can work around, must keep tuning and adding hacks, e.g.
 - Increasing # of threads & reducing call stack
 - Prepare thread pools for each shard
- Shall we just go asynchronous, please?
 - Less tuning points
 - Memory size & # of event loops
 - Better resource utilization with concurrent calls + less threads

Problems with large payloads

- We solved blocking problem with asynchronous programming, but can we send *10MB* personalized response to *100K* clients?
 - Can't hold that much in RAM – $10\text{MB} \times 100\text{K} = 1\text{TB}$
- What if we · they send too fast?
 - Different bandwidth & processing power
- We need '*just enough* buffering.'
 - Expect `OutOfMemoryError` otherwise.



Traditional vs. Reactive



@armeria_project



online/armeria

Reactive HTTP/2 proxy in 6 lines

```
// Use Armeria's async & reactive HTTP/2 client.  
HttpClient client = HttpClient.of("h2c://backend");  
Server server = Server.builder()  
    .http(8080)  
    .service("prefix:/",  
            // Forward all requests reactively.  
            (ctx, req) -> client.execute(req))  
    .build();
```

- Full example:

<https://github.com/line/armeria-examples/tree/master/proxy-server>



@armeria_project



Line/armeria

1st-class RPC support with better-than-upstream experience



@armeria_project



online/armeria

RPC vs. HTTP impedance mismatch

- RPC has been hardly a 1st-class citizen in web frameworks.
 - Which method was called with what parameters?
 - What's the return value? Did it succeed?

```
POST /some_service HTTP/1.1
Host: example.com
Content-Length: 96

<binary request>
```

```
HTTP/1.1 200 OK
Host: example.com
Content-Length: 192

<binary response>
```

Failed RPC call

```
192.167.1.2 - - [10/Oct/2000:13:55:36 -0700]
"POST /some_service HTTP/1.1" 200 2326
```



@armeria_project



online/armeria

Killing many birds with Structured Logging

- Timings
 - Low-level timings, e.g. DNS · Socket
 - Request · Response time
- Application-level
 - Custom attributes
 - User
 - Client type
 - Region, ...
- HTTP-level
 - Request · Response headers
 - Content preview, e.g. first 64 bytes
- RPC-level
 - Service type
 - method and parameters
 - Return values and exceptions



@armeria_project



Qline/armeria

First things first – Decorators

```
GrpcService.builder().addService(new MyServiceImpl()).build()
    .decorate((delegate, ctx, req) -> {
        ctx.log().addListener(log -> {
            ...
        }, RequestLogAvailability.COMPLETE);

        return delegate.serve(ctx, req);
    });
}
```

- Decorators are used everywhere in  Armeria
 - Most features mentioned in this presentation are decorators.



Async retrieval of structured logs

```
GrpcService.builder().addService(new MyServiceImpl()).build()
    .decorate(delegate, ctx, req) -> {
    ctx.log().addListener(log -> {
        ...
    }, RequestLogAvailability.COMPLETE);

    return delegate.serve(ctx, req);
});
```

Async retrieval of structured logs (cont'd)

```
ctx.log().addListener(log -> {
    long reqStartTime = log.requestStartTimeMillis();
    long resStartTime = log.responseStartTimeMillis();

    RpcRequest rpcReq = (RpcRequest) log.requestContent();
    if (rpcReq != null) {
        String method = rpcReq.method();
        List<Object> params = rpcReq.params();

        RpcResponse rpcRes = (RpcResponse) log.responseContent();
        if (rpcRes != null) {
            Object result = rpcRes.getNow(null);
        }
    }
}, RequestLogAvailability.COMPLETE);
```



```
t app.type          ✎ ✎ ✎ * ANDROID
t app.version       ✎ ✎ ✎ * 9.15.0
t args_json         ✎ ✎ ✎ * {"request":{...}}
t client_ip         ✎ ✎ ✎ *
t exception         ✎ ✎ ✎ *
t method             ✎ ✎ ✎ * SquareService#fetchMyEvents
t phase              ✎ ✎ ✎ * RELEASE
# processing_time_millis ✎ ✎ ✎ * 3
t request_header.name ✎ ✎ ✎ * fetchMyEvents
# request_header.seqid ✎ ✎ ✎ * 1
t request_header.type ✎ ✎ ✎ * CALL
t request_id         ✎ ✎ ✎ * 15710383415390
⌚ request_timestamp   ✎ ✎ ✎ * October 14th 2019, 16:32:21.539
t response_header.name ✎ ✎ ✎ * fetchMyEvents
# response_header.seqid ✎ ✎ ✎ * 1
t response_header.type ✎ ✎ ✎ * REPLY
t result_json        ✎ ✎ ✎ * {"success":{...}, "error":{}}
t server_id          ✎ ✎ ✎ *
# user.id            ✎ ✎ ✎ *
t user.user_region    ✎ ✎ ✎ * TH
```



Armeria



elasticsearch



kibana

Making a debug call

- Sending an ad-hoc query in RPC is hard.
 - Find a proper service definition, e.g. .thrift or .proto files
 - Set up code generator, build, IDE, etc.
 - Write some code that makes an RPC call.
- HTTP in contrast:
 - cURL, telnet command, web-based tools and more.
- What if we build something more *convenient* and *collaborative*?



Armeria documentation service

- Enabled by adding DocService
- Browse and invoke RPC services in an  Armeria server
 - No fiddling with binary payloads
 - Send a request without writing code
- Supports gRPC, Thrift and annotated services
- We have a plan to add:
 - Metric monitoring console
 - Runtime configuration editor, e.g. logger level



@armeria_project



Online/armeria

Services

Cassandra

POST	add()
POST	batch_mutate()
POST	describe_cluster_name()
POST	describe_keyspace()
POST	describe_keyspaces()
POST	describe_partitioner()
POST	describe_ring()
POST	describe_schema_versions()
POST	describe_snitch()
POST	describe_splits()
POST	describe_version()
POST	execute_cql_query()
POST	execute_prepared_cql_query()
POST	get()
POST	get_count()
POST	get_indexed_slices()
POST	get_range_slices()
POST	get_slice()
POST	insert()

Cassandra.add()

Increment or decrement a counter.

Parameters

Name	Required	Type	Description
key	required	binary	
column_parent	required	ColumnParent	
column	required	CounterColumn	
consistency_level	required	ConsistencyLevel	

Return Type

void

Exceptions

[InvalidRequestException](#)

[TimedOutException](#)

[UnavailableException](#)

- Share the URL to reproduce a call.

ⓘ ost:3000/docs/#/methods/com.linecorp.armeria.service.test.thrift.main.HelloService/hello/POST?request_body={"name": "world!"} ✓ ⋮ 🌐 ⚙ ☆

Debug

HTTP HEADERS

REQUEST BODY

```
{  
  "name": "world!"}  
}
```

```
{  
  "method": "hello",  
  "type": "REPLY",  
  "seqid": 0,  
  "args": {  
    "success": "Hello world!"  
  }  
}
```

SUBMIT **COPY AS A CURL COMMAND**

Cool features not available in upstream

- gRPC
 - Works on both HTTP/1 and 2
 - gRPC-Web support, i.e. can call gRPC services from JavaScript frontends
- Thrift
 - HTTP/2, TTEXT (human-readable REST-ish JSON)
- Can leverage  Armeria decorators
 - Structured logging, Metric collection, Distributed tracing, Authentication
 - CORS, SAML, Request throttling, Circuit breakers, Automatic retries, ...



@armeria_project



online/armeria

Cool features not available in upstream

- Can mix gRPC, Thrift, REST, Tomcat, Jetty, ...
 - on a single HTTP port & single JVM
 - without any proxies
 - REST API
 - Exposing metrics
 - Traditional JEE webapps
 - Static files
 - Health-check requests from load balancers
- Share common logic between different endpoints!

Unopinionated integration & migration



@armeria_project



online/armeria

Armeria 😍 What You 😍

- Use your favorite tech, not ours:
 - DI –  **spring**®, Guice, Dagger, ...
 - Protocols –  gRPC, Thrift, REST, ...
- Choose only what you want:
 - Most features are optional.
 - Compose and customize at your will.
 - Your application grows with you, not by its own.



@armeria_project



online/armeria

Case of slack

- Using Thrift since 2015
- Migrated from Thrift to gRPC
 - Can run both while clients are switching
- Leverages built-in non-RPC services:
 - PrometheusExpositionService 
 - HealthCheckService
 - BraveService – Distributed tracing with  honeycomb
 - DocService



@armeria_project



online/armeria

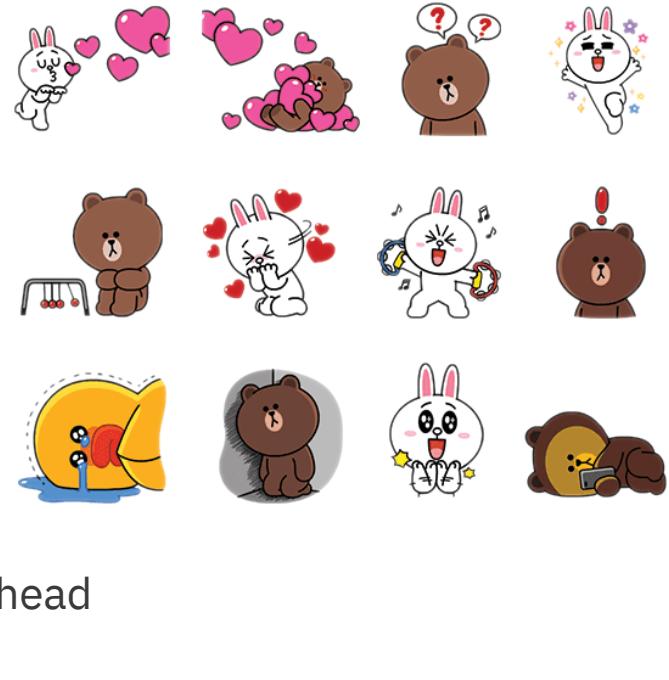
```
.defaultRequestTimeoutMillis(config.getRequestTimeoutMillis())
.maxNumConnections(config.getMaxConnections())
.meterRegistry(config.getMetricsRegistry())
.port(config.getPort(), SessionProtocol.HTTP)
.serviceUnder(config.getHealthPath(), healthCheckService)
.serviceUnder(config.getMetricsPath(), metricsService)
.serviceUnder(config.getDocsPath(), docService);
```

```
// Add user defined services.
config.getRawServices().forEach((path, service) -> builder.serviceUnder(path, service));
config.getThriftServices().forEach((path, service) -> builder.serviceUnder(path, service));
if (!config.getGrpcServices().isEmpty()) {
    GrpcServiceBuilder grpcBuilder = new GrpcServiceBuilder();
    config.getGrpcServices().forEach(service -> grpcBuilder.addService(service));
    builder.service(grpcBuilder.build());
}
```

- Full migration story: <https://sched.co/L715>

Case of LINE

- In-app emoji · sticker store (50k-100k reqs/sec)
- Before:
 - Spring Boot + Tomcat (HTTP/1) + Thrift on Servlet
 - Apache HttpClient
- After – Migrate keeping what you love 🤘
 - Spring Boot +  Armeria (HTTP/2)
 - Keep using Tomcat via TomcatService for the legacy
 - Thrift served directly & asynchronously = No Tomcat overhead
 - Armeria's HTTP/2 client w/ load-balancing

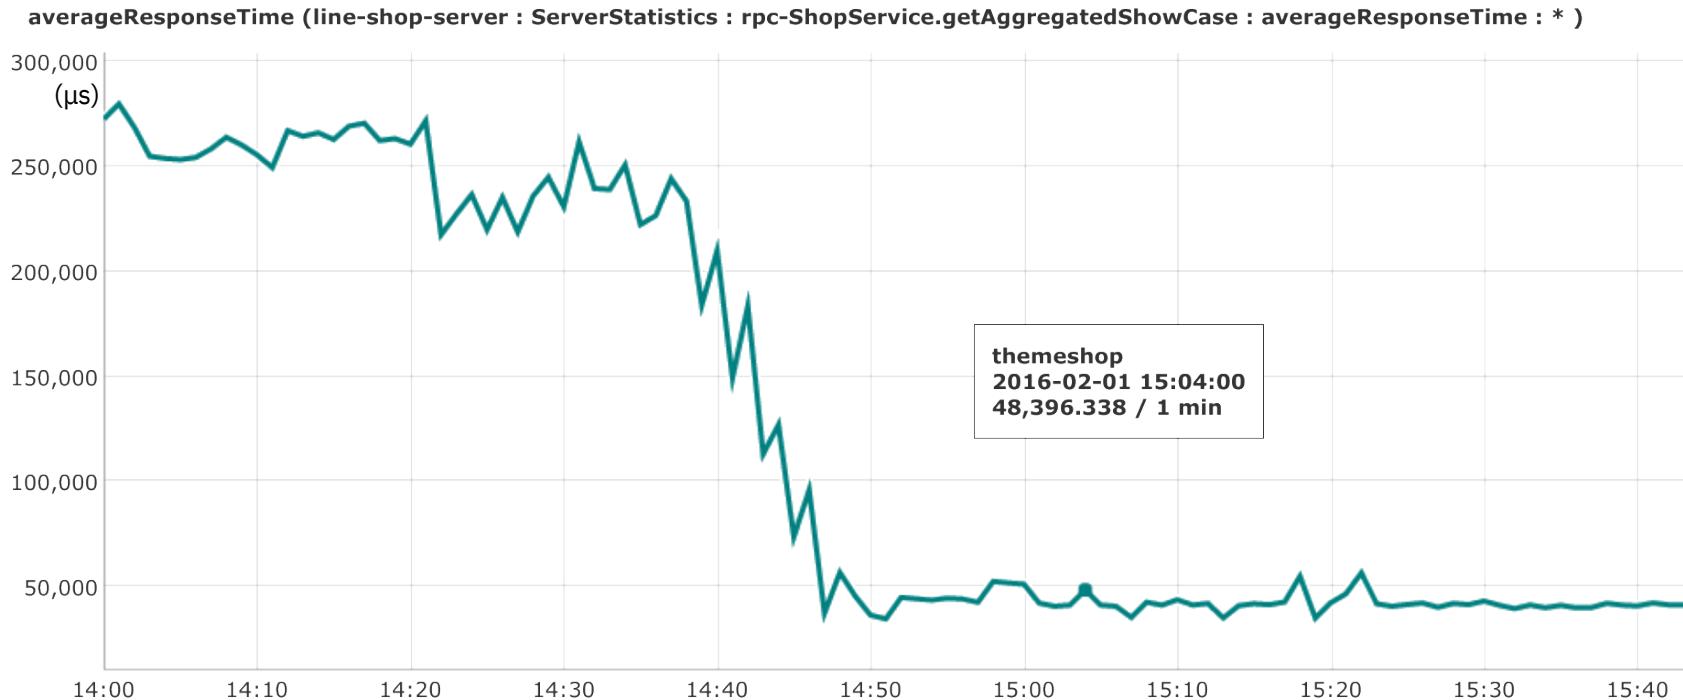


@armeria_project



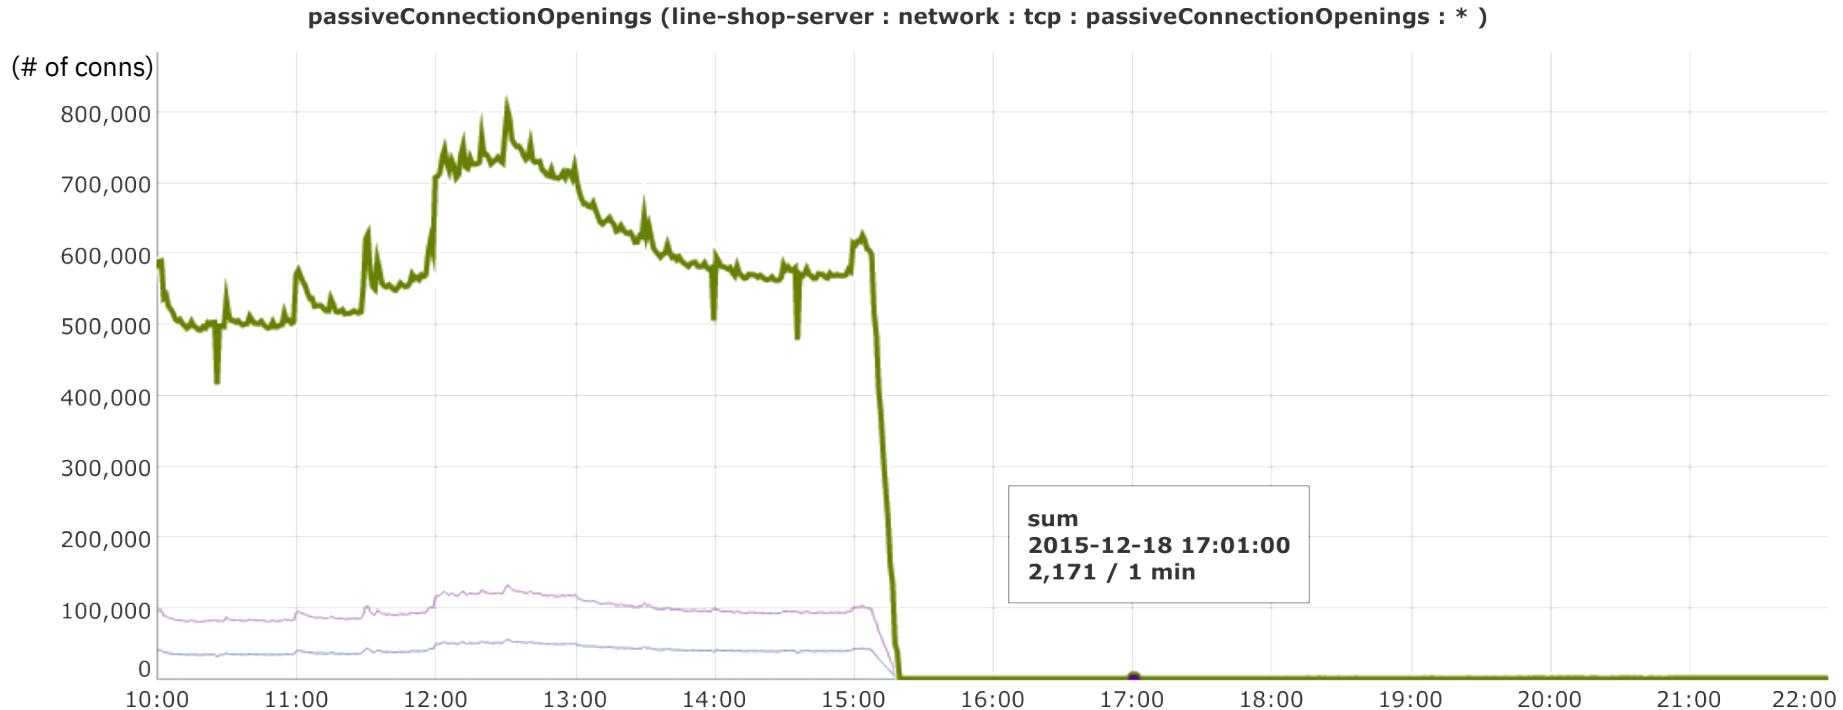
line/armeria

Case of LINE



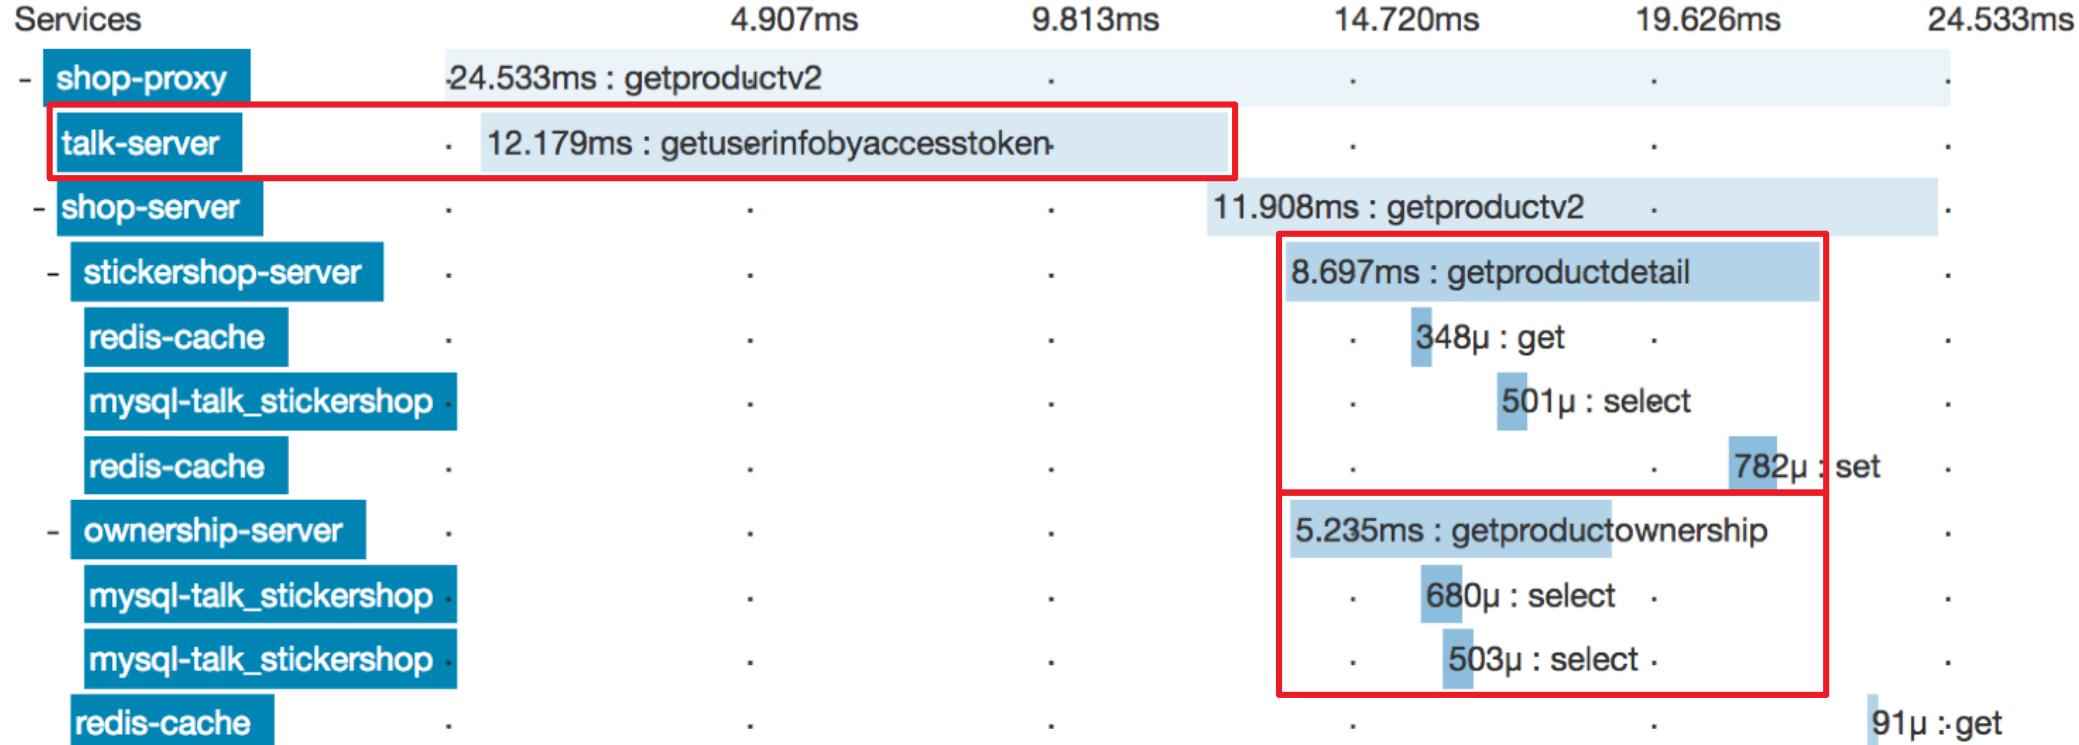
- Asynchronization of 3 synchronous calls

Case of LINE



- Significant reduction of inter-service connections

Case of LINE



- Distributed tracing with  ZIPKIN by just adding BraveService
- Full story: <https://www.slideshare.net/linecorp/line-zipkin>

Case of kakaopay

- Firm banking gateway
 - Talking to Korean banks via VAN (value-added network)
-  **Kotlin** +  Armeria
 - Mostly non-null API
 - Using @Nullable annotation extensibly
- Spring WebFlux + gRPC
- Armeria Replaces Spring's network layer (reactor-netty)
- gRPC served directly = No WebFlux overhead



@armeria_project



online/armeria

Less points of failure

Client-side load-balancing

Load balancers • Reverse proxies

- Pros
 - Distributes load
 - Offloads TLS overhead
 - Automatic health checks
 - Service discovery (?)
- Cons
 - More points of failure
 - Increased hops · latency
 - Uneven load distribution
 - Cost of operation
 - Health check lags



Client-side load balancing

- Client-side load balancing
 - Chooses endpoints *autonomously*
 - Service discovery – DNS,  **kubernetes**,  APACHE ZooKeeper™, ...
 - Near real-time health checks
 - Less points of failure
- Proxy-less Armeria server
 - OpenSSL-based high-performance TLS
 -  Netty + /dev/epoll
 - Assemble your services into a single port + single JVM!

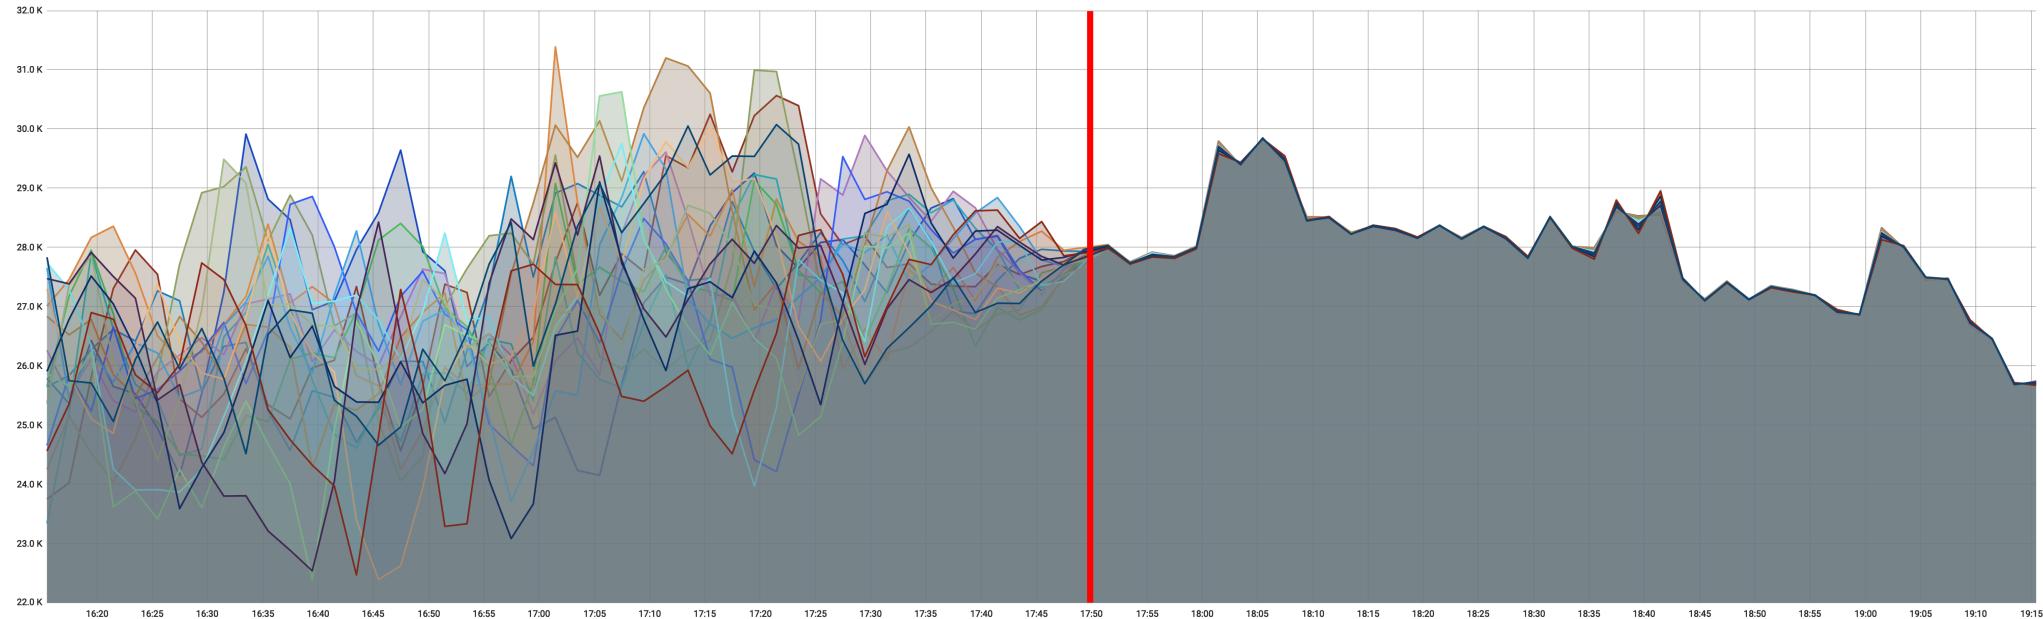


@armeria_project



online/armeria

HTTP/2 load distribution at LINE



- Full migration story:

https://speakerdeck.com/line_developers/lesson-learned-from-the-adoption-of-armeria-to-lines-authentication-system



@armeria_project



line/armeria

Near real-time health check

- Leverage HTTP/2 + long-polling
 - Significantly reduced number of health check requests, e.g. every 10s vs. 5m
 - Immediate notification of health status
- Server considered unhealthy
 - On disconnection
 - On server notification, e.g. graceful shutdown, self-test failure
- Fully backwards-compatible
 - Activated only when server responds with a special header



@armeria_project



Online/armeria

Client-side load-balancing with auto-retry and circuit breaker in 8 lines

```
// Kubernetes-style service discovery + long polling health check
EndpointGroup group = HealthCheckedEndpointGroup.of(
    DnsServiceEndpointGroup.of("my-service.cluster.local"),
    "/internal/healthcheck");
// Register the group into the registry.
EndpointGroupRegistry.register("myService", group, WEIGHTED_ROUND_ROBIN);
// Create an HTTP client with auto-retry and circuit breaker.
HttpClient client = HttpClient.builder("http://group:myService")
    .decorator(RetryingHttpClient.newDecorator(onServerErrorStatus()))
    .decorator(CircuitBreakerHttpClient.newDecorator(...))
    .build();
// Send a request.
HttpResponse res = client.get("/hello/armeria");
```



Future work

Consider joining us!



@armeria_project



online/armeria

The road to 1.0

(and beyond)

- Currently at 0.95
- Hoping to release before the end of 2019
- API stabilization · clean-up
- Post-1.0
 - Kotlin · Scala DSL
 - Evolving DocService to DashboardService
 - More transports & protocols
 - Web Sockets, UNIX domain sockets, Netty handlers, ...
- More decorators
- More service discovery mechanisms
 - Eureka, Consul, etcd, ...
- OpenAPI spec (.yml) generator
- Performance optimization



@armeria_project



online/armeria

Meet us at GitHub



github.com/line/armeria

line.github.io/armeria



@armeria_project



line/armeria